



Parallel Computing in R

from `parallel` to `foreach` and `future`

Chao Cheng

School of Statistics and Management
Shanghai University of Finance and Economics, China

November 22, 2022





Outline

- 1 R with C/C++
- 2 Parallel Computation in R
 - Introduction
 - The Parallel package
 - Advanced topics
 - Easy parallel computation
 - Other stuff
- 3 Acknowledgement





Extention of Rcpp

RcppProgress

- Article: Using RcppProgress to control the long computations in C++

```
# Rcpp::sourceCpp("rcpp_progress.cpp")
# long_computation(3000)
# 0%    10    20    30    40    50    60    70    80    90    100%
# [----|----|----|----|----|----|----|----|----|----|
# *****|
# [1] 3002.32
```





Extention of Rcpp

RcppParallel

```
1 #include<Rcpp.h>
2 // [[Rcpp::depends(RcppParallel)]]
3 #include <RcppParallel.h>
4
5 struct My_Sum : RcppParallel::Worker{
6     // member variables
7     const RcppParallel::RVector<double> input;
8     double res;
9
10    // constructor
11    My_Sum(Rcpp::NumericVector x) : input(x), res(0.0) {}
12    My_Sum(const My_Sum& my_sum, RcppParallel::Split): input(my_sum.input), res(0.0) {}
13    // operator functions
14    void operator()(std::size_t start, std::size_t end){
15        res += std::accumulate(input.begin() + start,
16                               input.begin() + end,
17                               0.0);
18    }
19    void join(const My_Sum& rhs){
20        res += rhs.res;
21    }
22 };
23 // [[Rcpp::export]]
24 double par_sum(Rcpp::NumericVector invec){
25     My_Sum my_sum(invec);
26     RcppParallel::parallelReduce(0, invec.size(), my_sum);
27     return(my_sum.res);
28 }
29 }
```

```
24
13:1 [Top Level]
Console Terminal Source Cpp Jobs
E:/[Doctor]/GroupSeminar/20191021_prepare/codes/
> library(rbenchmark)
> library(Rcpp)
> sourceCpp("rcpp_par2.cpp")
> x <- as.numeric(1 : 1000000)
> res1 <- sum(x)
> res2 <- par_sum(x)
> res3 <- rcpp_sum(x)
> identical(res1, res2)
[1] TRUE
> identical(res2, res3)
[1] TRUE
> benchmark(sum(x), par_sum(x), rcpp_sum(x), order = "relative")[, 1 : 4]
      test replications elapsed relative
2 par_sum(x)      100      0.03      1.000
1  sum(x)         100      0.11      3.667
3 rcpp_sum(x)     100      0.11      3.667
> |

of all subranges to which

as for parallel_for.

splitting constructor does two
```



Introduction

A toy example

- $1 + 2 + \dots + 100$
- Compute the sum of a vector

Abstraction

- The whole job can be break into small parts and they can be done independently of each other.
- Map + Reduce

Useful cases

- Simulation
- Bootstrap, MCMC and cross validation, etc.
- Elementwisely update an vector in ADMM algorithm (Parallel in C++)



Basic parallel computation for simulation

- Start multiple R sessions
- Preparation: load necessary packages, etc.
- Run simulation scripts, possibly according to session ID.
- Collect and summary the results by hand.

Abstraction

- Create workers
- Prepare workers
- Run script in parallel and collect the results.





The parallel package

- It's derived from `snow` and `multicore` packages.
- Useful reference:
 - Parallel R. (This book is a bit old.)
 - `parallel`'s documentations.
 - `parallel`'s vignettes.





A simple template

```
library(parallel)
# use all the cores of this machine
cls <- makeCluster(detectCores())

# initializing workers
clusterEvalQ(cls, fun)
# pass VARLIST from master to all the workers
clusterExport(cls, VARLIST)

# split full index of all tasks to workers
idx_split <- clusterSplit(cls, idx_full)
# carry out the task parFUNCTION parallely
res <- parLapply(cls, idx_split, parFUNCTION)

# stop workers
stopCluster(cls)
```




A simple example

```
library(parallel)
a <- rnorm(12)
slow_function <- function(invec){
  ...      # a slow function
}

cls <- makeCluster(4)
ind_seq <- clusterSplit(cls, a)
clusterExport(cls, varlist = "slow_function")
res_par <- parSapply(cls, ind_seq, slow_function)
res <- sum(res_par)
```





PSOCK vs FORK

```
a <- rnorm(100)
```

PSOCK

```
cls <- makeCluster(4)      # default type is PSOCK
```

FORK

```
cls <- makeCluster(4, type = "FORK")  # NOT available on Windows
```

```
parSapply(cls, 1 : 10, function(id){  
  return(a[id])  
})
```





PSOCK vs FORK

PSOCK

- Pros:
 - Use socket connection, a general approach.
 - All system, locally or remotely with suitable setup such as MPI
- Cons:
 - Might be hard to configure.
 - Manually transport the data.

FORK

- Pros: use FORK mechanism, no worry about variable transportation.
- Cons: Only for one machine, not available on Windows.





Parallel random number generation

- Manually use `set.seed()` on every worker.
- Use 'L'Ecuyer-CMRG' multiple RNG stream.
 - 1 `RNGkind("L'Ecuyer-CMRG")` on your main session.
 - 2 `set.seed()` on your main session.
 - 3 `clusterSetupRNGstream()` to set your workers' seed.





Variable transportation

- Explicit functions and variables will always be transported.
- FORK will copy the main session at creation.
- Others should be taken care of by hand.
- Additional configuration of `clusterExport` when nested in a function call.





Dark time of parallel computation

There are so many different parallel backends:

- snow
- multicore
- parallel
- MPI
- Redis
- Hadoop
- Spark
- Slurm
- ...

How to support them? How to maintain code?





Foreach

`foreach` defines a simple but powerful framework for map/reduce parallel computation.

Package author/code writer

Decide **which part of code** can run in parallel.

End user

Decide **how** to run in parallel based on their available resources.

`foreach` is syntactically structured in the form of a for loop.





Foreach

```
library(foreach)
# library(doParallel)
# registerDoParallel()
a <- 10
foreach(i = 1 : 12, j = 12 : 1, .combine = rbind) %dopar%{
  Sys.sleep(0.5)
  # print will be dropped when run in parallel
  print(paste("i = ", i, ", j = ", j, sep = " "))
  data.frame(i, j, a)
}
```





future and future.apply

- `future` provides a simple and uniform way of evaluating R expressions asynchronously using various resources available to the user.
- `future.apply` provides worry-free parallel alternatives to base-R "apply" functions.

```
library(future.apply)      # default plan is sequential
# plan(cluster)
x <- rnorm(16)
future_lapply(1 : 5, function(id){
  print(paste("id = ", id, sep = " "))      # normal print kept
  Sys.sleep(0.5)
  sum(x[1 : id])
})
```





- Asynchronous computation. Not constrained by a for-loop or apply syntax.
- Available extensions:
 - `future.apply`
 - `doFuture`: **backends for** `foreach`, `BiocParallel` **and** `plyr`.
 - `furrr`





future

```
library(future)
plan(cluster)

x <- future({
  x <- matrix(rnorm(10 ^ 6), nrow = 10 ^ 3)
  for(i in 1 : 5){
    print(paste("i = ", i))
    res <- eigen(x)
  }
  return(res)
}, seed = T)      # not block the main session

resolved(x)      # check whether the future is resolved
a <- rnorm(10)    # we can do other stuff at the main session
```





Personal suggestions

- Nested parallel is **NOT** recommended. At least it should be done with careful configuration.
- `future.apply` **VS** `foreach`
 - Familiar with `foreach`: just use the `doFuture` backends.
 - New to parallel: `future.apply` is a good start point for your code.
 - `future` backend will relay the printed messages.
 - Performance in parallel are close, **so-called**.
 - Performance for sequential are slower than for-loop.





I want progress bars

- `RcppProgress` allows to display a progress bar in the R console for long running computations taking place in c++ code, supports OpenMP.
- `pbapply` is a lightweight package that adds progress bar to vectorized R functions ('*apply'). It supports several parallel backends.
- `progress` shows ASCII progress bars.
- `progressr` provides a minimal API for reporting progress updates in R.
 - Developer is responsible for providing progress updates.
 - End user decides if, when, and how progress should be presented.





progressr

```
library(progressr)

slow_sum <- function(x) {
  p <- progressr::progressor(along = x)
  sum <- 0
  for (kk in seq_along(x)) {
    Sys.sleep(0.5)
    sum <- sum + x[kk]
    p(message = sprintf("Added %g", x[kk]))
  }
  sum
}

# handlers("default")      # default handler is "txtprogressbar"
with_progress(y <- slow_sum(1:10))
handlers("progress")
with_progress(y <- slow_sum(1:10))
```





Profile the future





Acknowledgement



Parallel computing with R
using foreach, future, and
other packages

BRYAN LEWIS



Future: Simple Async, Parallel
& Distributed Processing in R -
What's Next?

HENRIK BENGTSOON

